

* For execution: code - in memory

But entire code in memory - error code

↓
loaded

- limit of physical memory

- more memory

- large DS

Solution:

execute partially loaded program

- less memory - runs faster (less I/O need)
- more programs at same time
- increase CPU utilization
- increase throughput
- no increase → response time.

* VM: separation of user logical memory from physical memory

- execute partially loaded program


* Virtual Address Space: logical view of how process is stored in memory

implemented in 2 ways:

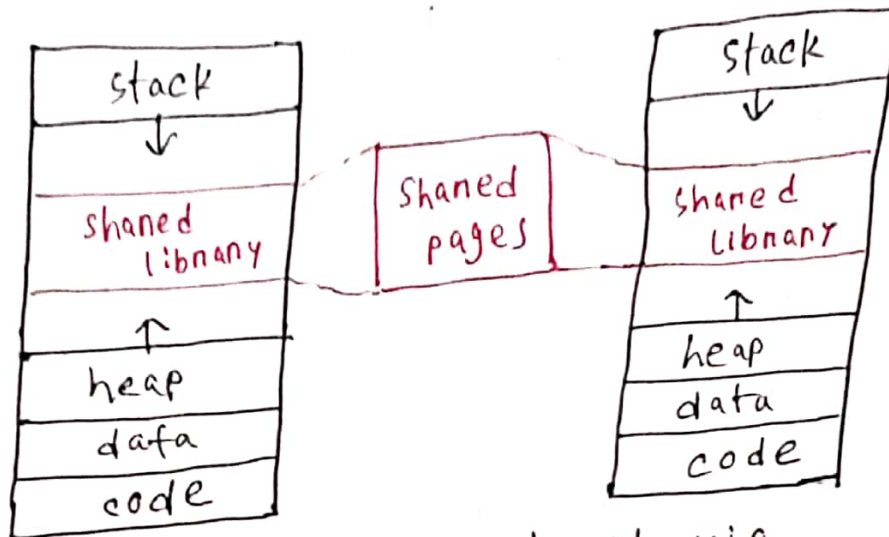
- 1) demand paging
- 2) demand segmentation

VM maps to PM ; $VM > PM$

↓
physical



* Shared library using VM:



(system library shared via mapping into virtual address space)

* Demand paging: need to execute an instruction from page

Less memory
Less I/O
fast response
more users

↓
demand for that page

→ load a page into main memory
↓
(instead of entire process)

Step: when page needed → reference to the page
3 cases:
1) invalid refer — abort (~~error~~) — terminate
2) page not in memory — bring it to mem
3) page in memory — use it

* Lazy swapper: never swaps a pages into memory unless page will be needed

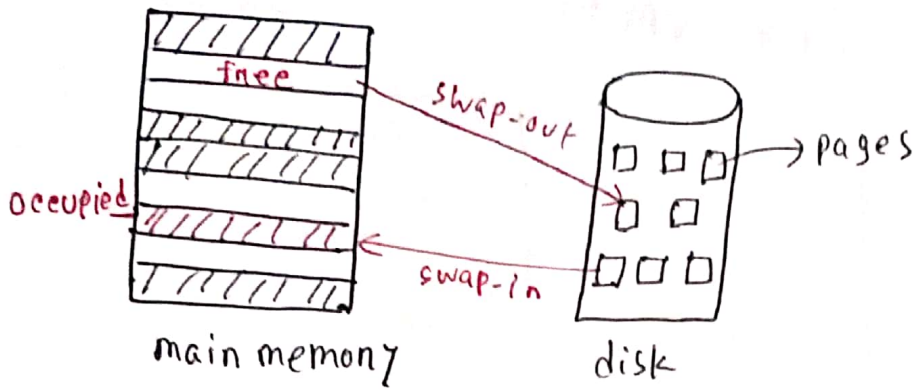
* Pagen: swapper that deal with pages

* Pure demand paging
* Locality of reference

Hardware needed for demand paging:

- 1) PTE with valid/invalid bit
- 2) secondary memory
- 3) instruction restart





Pros:

less memory
less swap time
increase degree of
multitasking

Cons: page fault

- Page not in memory

Fig: Lazy swapper (swap pages when needed)


* Valid - invalid bit:

↓
v - in memory

↓ i - not in memory
causes page fault

Initially, all bits - invalid (?)

↳ in page table entry (PTE)



* Steps in handling page fault:

- 1) reference to a page
- 2) page not in memory → set ⁰ bits in PTE
- 3) look for ^{missing} page in the backing store (disk)
- 4) bring it to free frame of main/physical memory.

5) reset PTE \rightarrow set bits 'v'

6) reset

- * No free frame in main memory
 - use page replacement algorithm
 - decrease page faults

Page replacement: Refers to a scenario in which a page from the main memory is replaced by a page from secondary memory

modify page
fault

- prevents over-allocation of memory
 - \downarrow
 - uses more memory than is physically available

modify
(dirty) bit

- reduce overhead of page transfers

determines which page is to be replaced.

Frame allocation algorithm

- determines \rightarrow 1) how many frame allocate
- 2) which frame \rightarrow replace



viset

* **FIFO Algo**: Replace page that comes in first

→ reference string of referenced page numbers:

7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 2, 1, 2, 0, 1, 7, 0, 1

→ 3 frames (3 pages in memory at a time per process)

7	0	1	2	0	3	0	4	2	3	0	3	2	1	2	0	1	7	0	1
7	7	7	2	2	2	2	4	4	4	0	0	0	0	0	0	0	7	7	7
	0	0	0	3	3	3	3	2	2	2	2	2	2	2	2	2	1	1	1
		1	1	1	1	0	0	0	3	3	3	3	3	3	3	3	2	2	2

15 page faults

→ more frame → more page faults

known as: **Belady's Anomaly**

* **Optimal Algo**: Replace page → not used for long period of time (in future)

7	0	1	2	0	3	0	4	2	3	0	3	2	1	2	0	1	7	0	1
7	7	7	2	2	2	2	4	4	4	0	0	0	0	0	0	0	7	7	7
	0	0	0	3	3	3	3	2	2	2	2	2	2	2	2	2	1	1	1
		1	1	1	1	0	0	0	3	3	3	3	3	3	3	3	2	2	2

→ optimal ans: 9 page faults

* **LRU (Least Recently Used) Algo**: Replace page — used earlier in past

7	0	1	2	0	3	0	4	2	3	0	3	2	1	2	0	1	7	0	1
7	7	7	2	2	2	2	4	4	4	0	0	0	0	0	0	0	7	7	7
	0	0	0	3	3	3	3	2	2	2	2	2	2	2	2	2	1	1	1
		1	1	1	1	0	0	0	3	3	3	3	3	3	3	3	2	2	2

→ 12 page faults — worse than optimal better than FIFO

* Counting Algo:

keep a count of
number of references
made to each page

1) LFU ^{Lease}

→ Frequently used

replace page
with
smallest count

2) MFU ^{most}

replace

page with smallest count that

- just आना हुआ (brought in)
- not used yet



3 major activities of demand paging:

- 1) Service the interrupt
- 2) Read the page
- 3) Restart the process

$$* EAT = (1-p) \times \text{memory access} + p (\text{page fault overhead} + \text{swap page out} + \text{swap page in})$$

Hence, EAT - Effective access time
1st part - hit, 2nd part - miss

p - page fault rate

$$0 \leq p \leq 1$$

$p=0$; no page fault
 $p=1$; every reference is a fault

$$\therefore 1-p = \text{hit rate}$$

a:

memory access time = 200 ns

Avg. page-fault service time = 8 ms = 8×10^6 ns

$$\therefore EAT = (1-p) \times 200 + p \times 8 \times 10^6$$
$$= 7,999,800p + 200$$

now, 1 access out of 1000 \rightarrow page fault

$$\therefore p = \frac{1}{1000} = 10^{-3}$$

$$\therefore EAT = 7,999,800 \times 10^{-3} + 200$$

$$= 8,199.8 \text{ ns}$$

$$= 8.2 \mu\text{s}$$

$$\text{slow down by factor of} = \frac{\text{miss}}{\text{hit}} = \frac{10^{-3} \times 8 \times 10^6}{(1 - 10^{-3}) \times 200} = 40$$

* if performance degradation $< 10\% = \frac{10}{100}$

$$\therefore \frac{\text{miss}}{\text{hit}} < \frac{10}{100}$$

$$\rightarrow \frac{8 \times 10^6 \times p}{(1 - p) \times 200} < \frac{1}{10}$$

$$\rightarrow 8 \times 10^6 \times p < 20 - 20p$$

$$\rightarrow p < \frac{20}{8 \times 10^6 + 20}$$

$$\rightarrow p < 2.5 \times 10^{-6}$$

$$\rightarrow p < \frac{1}{0.4 \times 10^6}$$

$$\rightarrow p < \frac{1}{4 \times 10^5}$$

less than
1 page fault in
4,00,000 mem. access